

```
from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
# Install necessary libraries  
!pip install tensorflow kagglehub opencv-python matplotlib scikit-learn pandas numpy Augmentor
```

Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (2.0.2)
Collecting Augmentor
 Downloading Augmentor-0.2.12-py2.py3-none-any.whl.metadata (1.3 kB)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)
Requirement already satisfied: gast!=0.5.0,!>0.5.1,!>0.5.2,>=0.2.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from tensorflow) (24.2)
Requirement already satisfied: protobuf!=4.21.0,!>4.21.1,!>4.21.2,!>4.21.3,!>4.21.4,!>4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/python3.1
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)
Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.0.1)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (4.13.2)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.71.0)
Requirement already satisfied: tensorboard<2.19,>=2.18 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.18.0)
Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.8.0)
Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.13.0)
Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.4.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.37.1)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.11/dist-packages (from kagglehub) (6.0.2)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from kagglehub) (4.67.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.2)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)

```
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow) (0.45.1)
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (13.9.4)
Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.0.8)
Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.15.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2025.1)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.8)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.1.3)
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tensorboard<2.19,>=2.18-
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (3.0.
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (2.
Requirement already satisfied: mdurl~0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->tenso
Downloading Augmentor-0.2.12-py2.py3-none-any.whl (38 kB)
Installing collected packages: Augmentor
Successfully installed Augmentor-0.2.12
```

```
!pip install ultralytics
```



```
-----  
Attempting uninstall: nvidia-curand-cu12  
  Found existing installation: nvidia-curand-cu12 10.3.6.82  
  Uninstalling nvidia-curand-cu12-10.3.6.82:  
    Successfully uninstalled nvidia-curand-cu12-10.3.6.82  
Attempting uninstall: nvidia-cufft-cu12  
  Found existing installation: nvidia-cufft-cu12 11.2.3.61  
  Uninstalling nvidia-cufft-cu12-11.2.3.61:  
    Successfully uninstalled nvidia-cufft-cu12-11.2.3.61  
Attempting uninstall: nvidia-cuda-runtime-cu12  
  Found existing installation: nvidia-cuda-runtime-cu12 12.5.82  
  Uninstalling nvidia-cuda-runtime-cu12-12.5.82:  
    Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82  
Attempting uninstall: nvidia-cuda-nvrtc-cu12  
  Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82  
  Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:  
    Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82  
Attempting uninstall: nvidia-cuda-cupti-cu12  
  Found existing installation: nvidia-cuda-cupti-cu12 12.5.82  
  Uninstalling nvidia-cuda-cupti-cu12-12.5.82:  
    Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82  
Attempting uninstall: nvidia-cublas-cu12  
  Found existing installation: nvidia-cublas-cu12 12.5.3.2  
  Uninstalling nvidia-cublas-cu12-12.5.3.2:  
    Successfully uninstalled nvidia-cublas-cu12-12.5.3.2  
Attempting uninstall: nvidia-cusparse-cu12  
  Found existing installation: nvidia-cusparse-cu12 12.5.1.3  
  Uninstalling nvidia-cusparse-cu12-12.5.1.3:  
    Successfully uninstalled nvidia-cusparse-cu12-12.5.1.3  
Attempting uninstall: nvidia-cudnn-cu12  
  Found existing installation: nvidia-cudnn-cu12 9.3.0.75  
  Uninstalling nvidia-cudnn-cu12-9.3.0.75:  
    Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75  
Attempting uninstall: nvidia-cusolver-cu12  
  Found existing installation: nvidia-cusolver-cu12 11.6.3.83  
  Uninstalling nvidia-cusolver-cu12-11.6.3.83:  
    Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83  
Successfully installed nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-cu12
```

```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
import os
```

```
import cv2
import shutil
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.applications import ResNet50
import kagglehub
import Augmentor
import pathlib
import glob
from IPython.display import display, Javascript
from ultralytics import YOLO
```

→ Creating new Ultralytics Settings v0.0.6 file ✓
View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.json'
Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see <https://docs.ultralytics.com/quickstart>

```
def keep_alive():
    display(Javascript('''
        function ClickConnect(){
            console.log("Working");
            document.querySelector("colab-toolbar-button#connect").click()
        }
        setInterval(ClickConnect, 60000)
    '''))
keep_alive()
```

```
print("Downloading HAM10000 dataset...")
path = kagglehub.dataset_download("kmader/skin-cancer-mnist-ham10000")
print("Path to dataset files:", path)
data_dir_train = pathlib.Path(path)
metadata_path = os.path.join(path, "HAM10000_metadata.csv")
images_path_part1 = os.path.join(path, "HAM10000_images_part_1")
images_path_part2 = os.path.join(path, "HAM10000_images_part_2")
```

→ Downloading HAM10000 dataset...
Path to dataset files: /kaggle/input/skin-cancer-mnist-ham10000

```
if not os.path.exists(metadata_path):
    raise FileNotFoundError("Metadata file not found.")
if not (os.path.exists(images_path_part1) and os.path.exists(images_path_part2)):
    raise FileNotFoundError("Images directory not found.")
metadata = pd.read_csv(metadata_path)

lesion_type_dict = {
    'nv': {'name': 'Melanocytic nevi', 'severity': 'Low'},
    'mel': {'name': 'Melanoma', 'severity': 'High'},
    'bkl': {'name': 'Benign keratosis-like lesions', 'severity': 'Moderate'},
    'bcc': {'name': 'Basal cell carcinoma', 'severity': 'High'},
    'akiec': {'name': 'Actinic keratoses', 'severity': 'Moderate'},
    'vasc': {'name': 'Vascular lesions', 'severity': 'Low'},
    'df': {'name': 'Dermatofibroma', 'severity': 'Low'}
}
metadata['cell_type'] = metadata['dx'].map(lambda x: lesion_type_dict[x]['name'])
metadata['severity'] = metadata['dx'].map(lambda x: lesion_type_dict[x]['severity'])
metadata['cell_type_idx'] = pd.Categorical(metadata['cell_type']).codes
class_names = list(lesion_type_dict.keys())

restructured_dir = "/content/ham10000_restructured"
os.makedirs(restructured_dir, exist_ok=True)
for split in ['train', 'val']:
    for class_name in class_names:
        os.makedirs(os.path.join(restructured_dir, split, 'images', class_name), exist_ok=True)
        os.makedirs(os.path.join(restructured_dir, split, 'labels', class_name), exist_ok=True)

train_meta, val_meta = train_test_split(metadata, test_size=0.2, random_state=123)
for _, row in train_meta.iterrows():
    image_id = row['image_id']
    class_label = row['dx']
    src_path = os.path.join(images_path_part1, f"{image_id}.jpg")
    if not os.path.exists(src_path):
        src_path = os.path.join(images_path_part2, f"{image_id}.jpg")
    dest_path = os.path.join(restructured_dir, 'train', 'images', class_label, f"{image_id}.jpg")
    if os.path.exists(src_path):
        shutil.copy(src_path, dest_path)
        with open(os.path.join(restructured_dir, 'train', 'labels', class_label, f"{image_id}.txt"), 'w') as f:
```

```
f.write(f"{class_names.index(class_label)} 0.5 0.5 1.0 1.0\n")

for _, row in val_meta.iterrows():
    image_id = row['image_id']
    class_label = row['dx']
    src_path = os.path.join(images_path_part1, f"{image_id}.jpg")
    if not os.path.exists(src_path):
        src_path = os.path.join(images_path_part2, f"{image_id}.jpg")
    dest_path = os.path.join(restructured_dir, 'val', 'images', class_label, f"{image_id}.jpg")
    if os.path.exists(src_path):
        shutil.copy(src_path, dest_path)
    with open(os.path.join(restructured_dir, 'val', 'labels', class_label, f"{image_id}.txt"), 'w') as f:
        f.write(f"{class_names.index(class_label)} 0.5 0.5 1.0 1.0\n")

for c in class_names:
    p = Augmentor.Pipeline(os.path.join(restructured_dir, 'train', 'images', c))
    p.rotate(probability=0.7, max_left_rotation=10, max_right_rotation=10)
    p.sample(500)
```

Initialised with 5385 image(s) found.
Output directory set to /content/ham10000_restructured/train/images/nv/output.Processing <PIL.JpegImagePlugin.JpegImageFile image mode=RGB size=600x450 at 0x7E
Initialised with 873 image(s) found.
Output directory set to /content/ham10000_restructured/train/images/mel/output.Processing <PIL.Image.Image image mode=RGB size=600x450 at 0x7E
Initialised with 874 image(s) found.
Output directory set to /content/ham10000_restructured/train/images/bkl/output.Processing <PIL.Image.Image image mode=RGB size=600x450 at 0x7E
Initialised with 414 image(s) found.
Output directory set to /content/ham10000_restructured/train/images/bcc/output.Processing <PIL.Image.Image image mode=RGB size=600x450 at 0x7E
Initialised with 263 image(s) found.
Output directory set to /content/ham10000_restructured/train/images/akiec/output.Processing <PIL.Image.Image image mode=RGB size=600x450 at 0x7E
Initialised with 112 image(s) found.
Output directory set to /content/ham10000_restructured/train/images/vasc/output.Processing <PIL.Image.Image image mode=RGB size=600x450 at 0x7E
Initialised with 91 image(s) found.
Output directory set to /content/ham10000_restructured/train/images/df/output.Processing <PIL.JpegImagePlugin.JpegImageFile image mode=RGB size=600x450 at 0x7E

```
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    os.path.join(restructured_dir, 'train', 'images'), batch_size=16, image_size=(128, 128),
    label_mode='categorical', seed=123, subset="training", validation_split=0.2
)
val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    os.path.join(restructured_dir, 'val', 'images'), batch_size=16, image_size=(128, 128),
```

```
label_mode='categorical', seed=123, subset="validation", validation_split=0.2
)

→ Found 11512 files belonging to 7 classes.
Using 9210 files for training.
Found 2003 files belonging to 7 classes.
Using 400 files for validation.
```

```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
```

```
def show_one_image_per_class(folder_path, class_names, title="Folder"):
    plt.figure(figsize=(16, 8))
    for i, class_name in enumerate(class_names):
        class_folder = os.path.join(folder_path, class_name)
        image_list = os.listdir(class_folder)
        if image_list:
            image_path = os.path.join(class_folder, image_list[0])
            img = mpimg.imread(image_path)
            plt.subplot(2, (len(class_names)+1)//2, i+1)
            plt.imshow(img)
            plt.title(class_name)
            plt.axis('off')
    plt.suptitle(f' One Image from Each Class in {title}', fontsize=16)
    plt.tight_layout()
    plt.show()

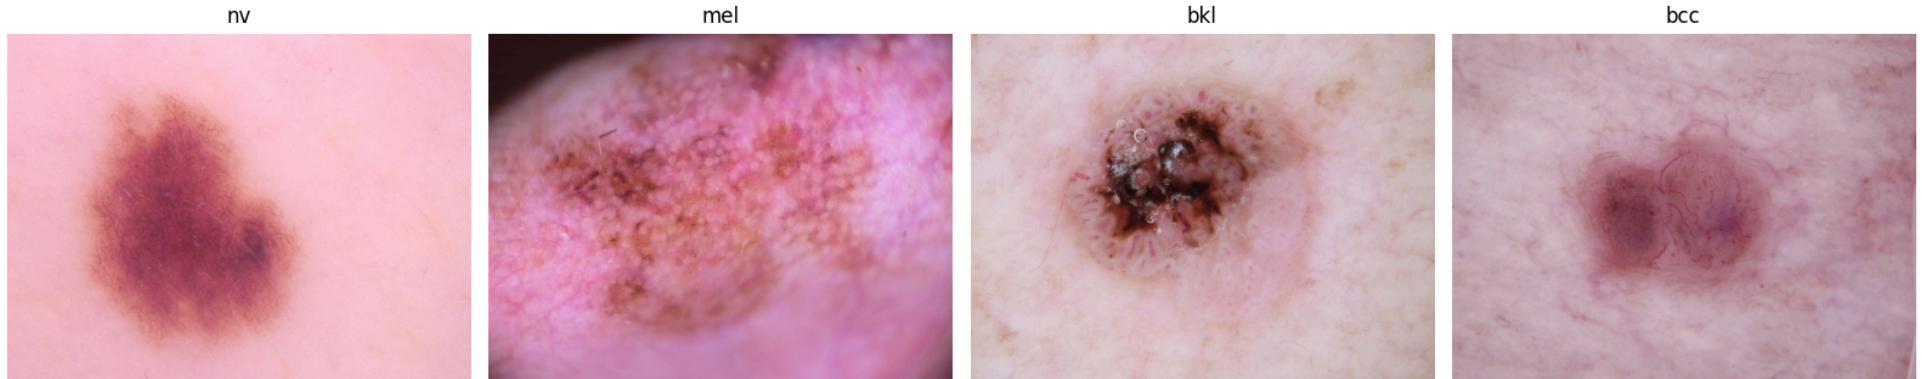
train_path = os.path.join(restructured_dir, 'train', 'images')
train_classes = class_names

val_path = os.path.join(restructured_dir, 'val', 'images')
val_classes = class_names

show_one_image_per_class(train_path, train_classes, title="Train")
show_one_image_per_class(val_path, val_classes, title="Validation")
```

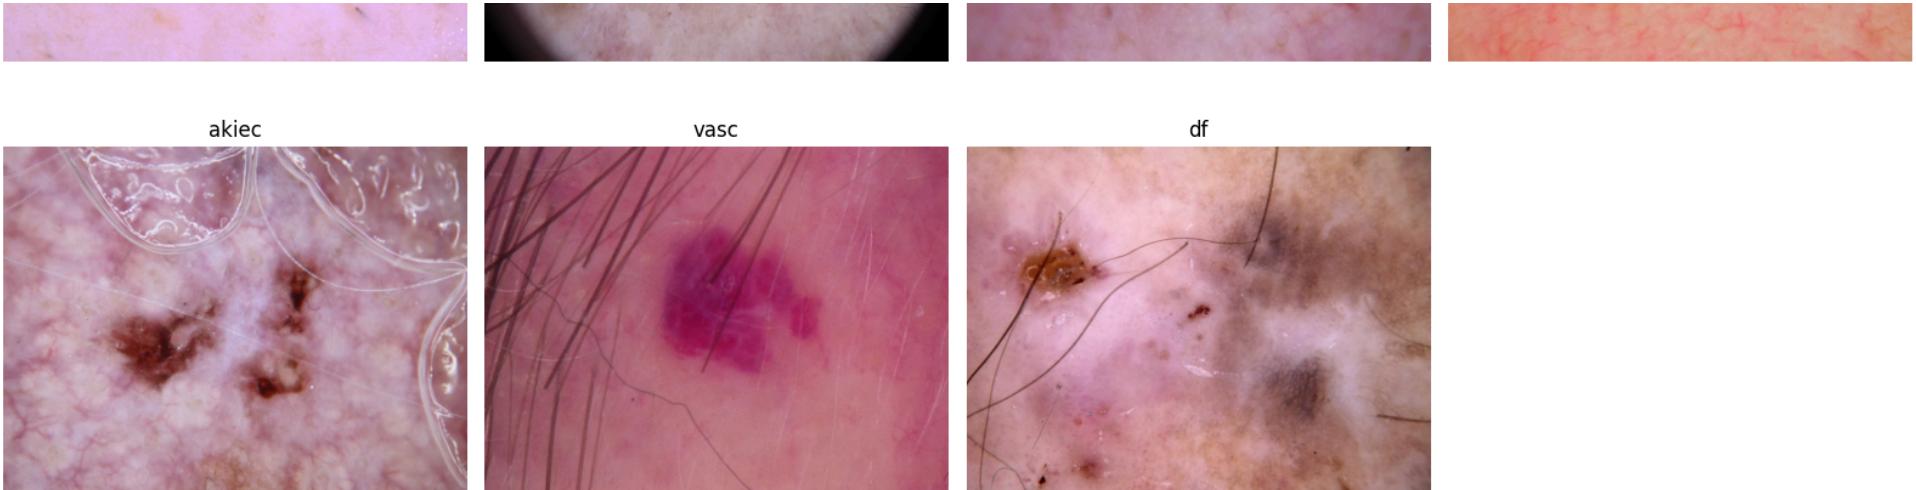
```
↳ <ipython-input-13-55b07c51f75f>:18: UserWarning: Glyph 128444 (\N{FRAME WITH PICTURE}) missing from font(s) DejaVu Sans.  
  plt.tight_layout()  
/usr/local/lib/python3.11/dist-packages/IPython/core/pylabtools.py:151: UserWarning: Glyph 128444 (\N{FRAME WITH PICTURE}) missing from fon  
  fig.canvas.print_figure(bytes_io, **kw)
```

□ One Image from Each Class in Train



□ One Image from Each Class in Validation





```
AUTOTUNE = tf.data.AUTOTUNE
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)

checkpoint_cnn = tf.keras.callbacks.ModelCheckpoint("cnn_model.h5", monitor="val_accuracy", save_best_only=True, verbose=1)
earlystop_cnn = tf.keras.callbacks.EarlyStopping(monitor="val_accuracy", patience=5, verbose=1)
checkpoint_resnet = tf.keras.callbacks.ModelCheckpoint("resnet_model.h5", monitor="val_accuracy", save_best_only=True, verbose=1)
earlystop_resnet = tf.keras.callbacks.EarlyStopping(monitor="val_accuracy", patience=5, verbose=1)

def create_cnn_model():
    model = models.Sequential([
        layers.Rescaling(1./255, input_shape=(128, 128, 3)),
        layers.Conv2D(32, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(128, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Dropout(0.5),
```

```
layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dropout(0.25),
layers.Dense(len(class_names), activation='softmax')
])
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
return model

import os
os.environ["CUDA_VISIBLE_DEVICES"] = "-1"

import os

if tf.config.list_physical_devices('GPU'):
    print("GPU found, using the GPU for training")
else:
    print("No GPU found, using the CPU for training")

cnn_model = create_cnn_model()
cnn_history = cnn_model.fit(train_ds, validation_data=val_ds, epochs=15, callbacks=[checkpoint_cnn, earlystop_cnn])

5/6/5/6 284s 494ms/step - accuracy: 0.5830 - loss: 1.0989 - val_accuracy: 0.6800 - val_loss: 0.8461
→ Epoch 3/15
576/576 0s 505ms/step - accuracy: 0.6381 - loss: 0.9626
Epoch 3: val_accuracy did not improve from 0.68000
576/576 332s 511ms/step - accuracy: 0.6381 - loss: 0.9626 - val_accuracy: 0.6775 - val_loss: 0.7966
Epoch 4/15
576/576 0s 488ms/step - accuracy: 0.6643 - loss: 0.8912
Epoch 4: val_accuracy did not improve from 0.68000
576/576 286s 497ms/step - accuracy: 0.6643 - loss: 0.8912 - val_accuracy: 0.6725 - val_loss: 0.8211
```

Epoch 8/15

576/576 0s 494ms/step - accuracy: 0.7694 - loss: 0.6062

Epoch 8: val_accuracy improved from 0.71250 to 0.71500, saving model to cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is consider

576/576 290s 504ms/step - accuracy: 0.7694 - loss: 0.6062 - val_accuracy: 0.7150 - val_loss: 0.7443

Epoch 9/15

576/576 0s 489ms/step - accuracy: 0.7919 - loss: 0.5497

Epoch 9: val_accuracy did not improve from 0.71500

576/576 318s 497ms/step - accuracy: 0.7919 - loss: 0.5497 - val_accuracy: 0.7125 - val_loss: 0.7966

Epoch 10/15

576/576 0s 489ms/step - accuracy: 0.8017 - loss: 0.5154

Epoch 10: val_accuracy improved from 0.71500 to 0.73500, saving model to cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is consider

576/576 284s 494ms/step - accuracy: 0.8017 - loss: 0.5154 - val_accuracy: 0.7350 - val_loss: 0.7340

Epoch 11/15

576/576 0s 500ms/step - accuracy: 0.8340 - loss: 0.4490

Epoch 11: val_accuracy did not improve from 0.73500

576/576 292s 506ms/step - accuracy: 0.8340 - loss: 0.4491 - val_accuracy: 0.7075 - val_loss: 0.7606

Epoch 12/15

576/576 0s 491ms/step - accuracy: 0.8309 - loss: 0.4458

Epoch 12: val_accuracy did not improve from 0.73500

576/576 317s 499ms/step - accuracy: 0.8309 - loss: 0.4458 - val_accuracy: 0.7175 - val_loss: 0.7915

Epoch 13/15

576/576 0s 488ms/step - accuracy: 0.8438 - loss: 0.4005

Epoch 13: val_accuracy did not improve from 0.73500

576/576 284s 493ms/step - accuracy: 0.8438 - loss: 0.4005 - val_accuracy: 0.6925 - val_loss: 0.8862

Epoch 14/15

576/576 0s 489ms/step - accuracy: 0.8538 - loss: 0.3963

Epoch 14: val_accuracy did not improve from 0.73500

576/576 285s 494ms/step - accuracy: 0.8538 - loss: 0.3963 - val_accuracy: 0.7175 - val_loss: 0.8357

Epoch 15/15

576/576 0s 490ms/step - accuracy: 0.8612 - loss: 0.3674

Epoch 15: val_accuracy did not improve from 0.73500

576/576 324s 498ms/step - accuracy: 0.8612 - loss: 0.3674 - val_accuracy: 0.7150 - val_loss: 0.8805

Epoch 15: early stopping

import matplotlib.pyplot as plt

train_acc = cnn_history.history['accuracy']
val_acc = cnn_history.history['val_accuracy']
epochs = range(1, len(train_acc) + 1)

plt.figure(figsize=(8, 5))

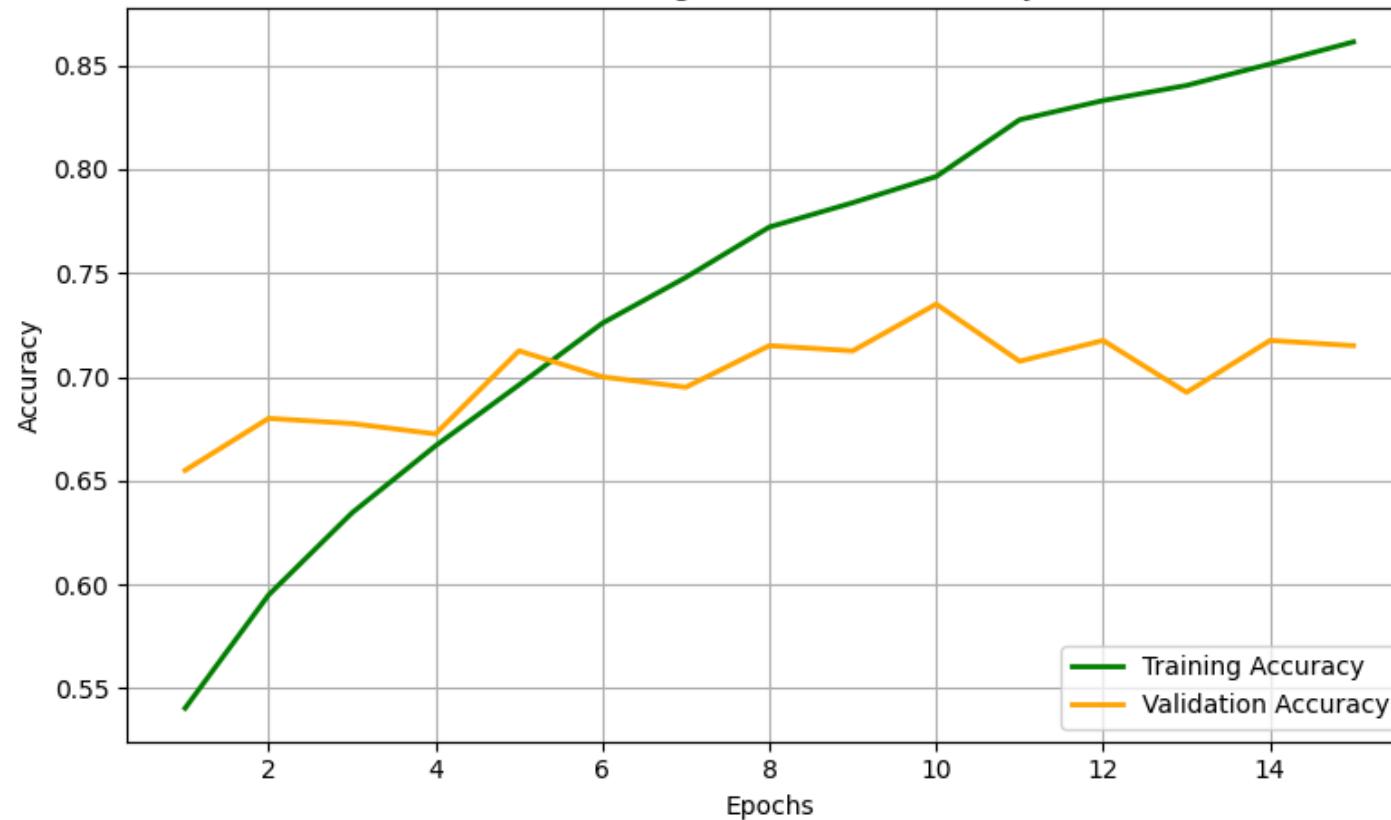
```
plt.plot(epochs, train_acc, 'g-', label='Training Accuracy', linewidth=2)
plt.plot(epochs, val_acc, 'orange', label='Validation Accuracy', linewidth=2)
plt.title('CNN Training vs Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

train_loss = cnn_history.history['loss']
val_loss = cnn_history.history['val_loss']

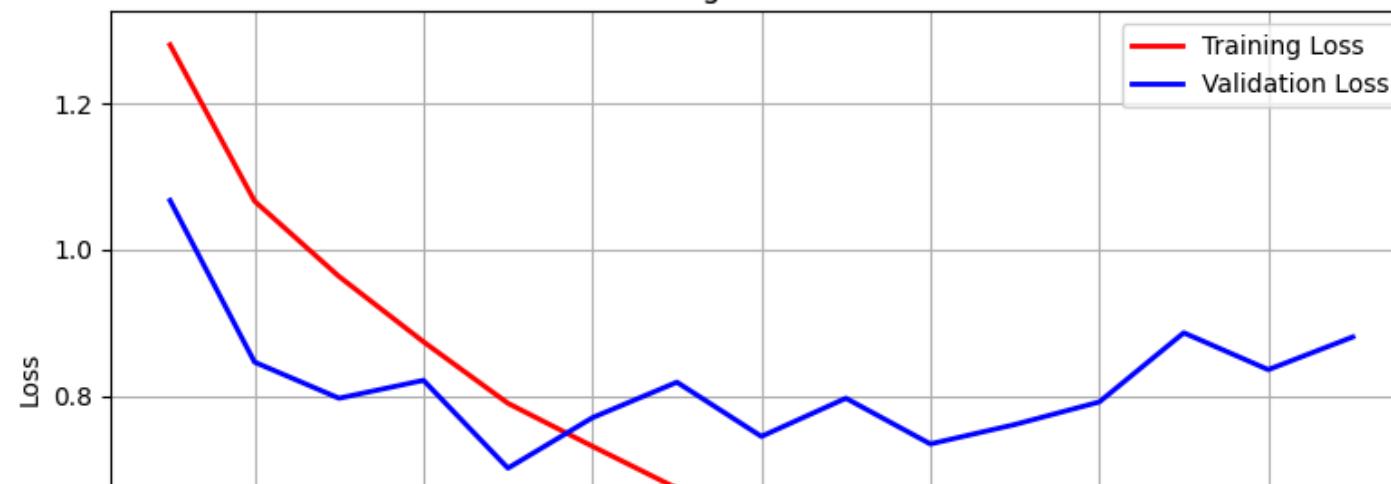
plt.figure(figsize=(8, 5))
plt.plot(epochs, train_loss, 'r-', label='Training Loss', linewidth=2)
plt.plot(epochs, val_loss, 'blue', label='Validation Loss', linewidth=2)
plt.title('CNN Training vs Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right')
plt.grid(True)
plt.tight_layout()
plt.show()
```

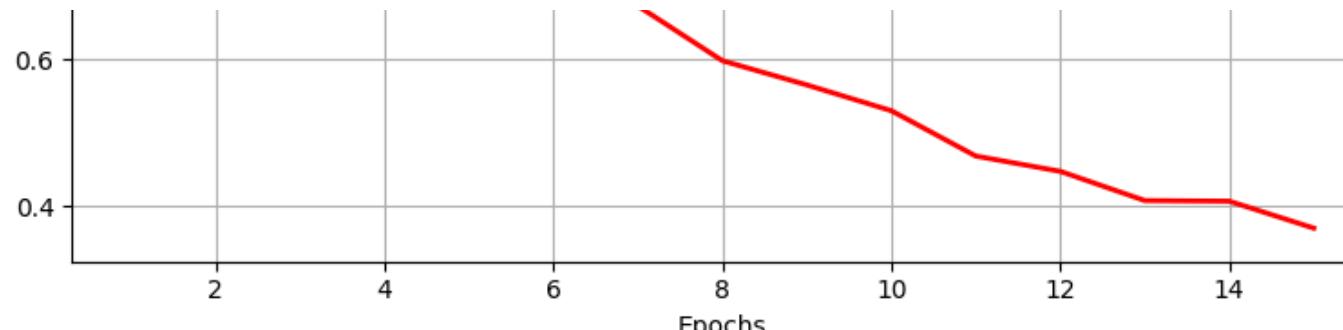


CNN Training vs Validation Accuracy



CNN Training vs Validation Loss





```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

test_ds = val_ds

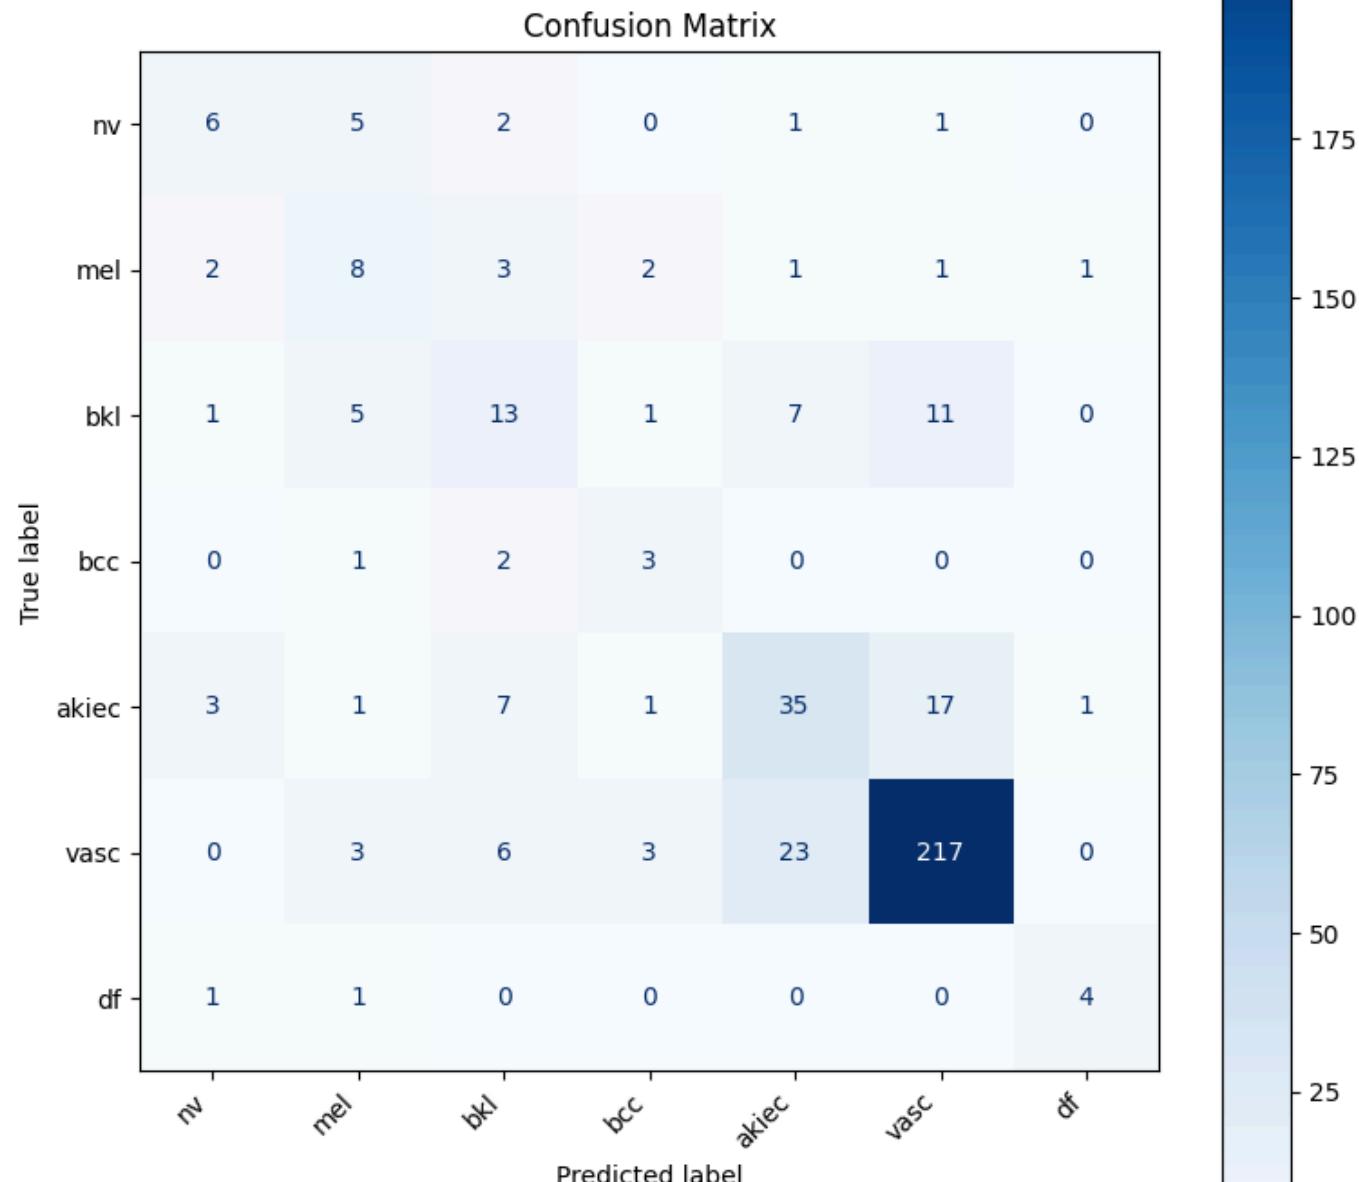
y_probs = cnn_model.predict(test_ds)
y_pred = np.argmax(y_probs, axis=1)

y_true_list = []
for X_batch, y_batch in test_ds:
    if y_batch.ndim > 1:
        y_true_list.append(np.argmax(y_batch, axis=1))
    else:
        y_true_list.append(y_batch)
y_true = np.concatenate(y_true_list, axis=0)

cm = confusion_matrix(y_true, y_pred)

disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                               display_labels=class_names)
fig, ax = plt.subplots(figsize=(8, 8))
disp.plot(ax=ax, cmap='Blues', colorbar=True)
plt.title('Confusion Matrix')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

25/25 3s 116ms/step



Start coding or generate with AI.

```
from tensorflow.keras.applications.resnet50 import preprocess_input, ResNet50

def create_resnet_model():

    base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(128, 128, 3))
    base_model.trainable = False

    model = models.Sequential([
        layers.Lambda(preprocess_input, input_shape=(128, 128, 3)),
        base_model,
        layers.GlobalAveragePooling2D(),
        layers.Dense(128, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(len(class_names), activation='softmax')
    ])

    model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    return model

resnet_model = create_resnet_model()
resnet_history = resnet_model.fit(train_ds, validation_data=val_ds, epochs=15, callbacks=[checkpoint_resnet, earlystop_resnet])
```

Download data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
94765736/94765736 0s 0us/step
Epoch 1/15
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/lambda_layer.py:65: UserWarning: Do not pass an `input_shape`/`input_dim` argument to `super().__init__(**kwargs)
576/576 0s 906ms/step - accuracy: 0.5415 - loss: 1.4260
Epoch 1: val_accuracy improved from -inf to 0.70500, saving model to resnet_model.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered
576/576 559s 954ms/step - accuracy: 0.5416 - loss: 1.4255 - val_accuracy: 0.7050 - val_loss: 0.7594

```
Epoch 2/15
576/576 0s 911ms/step - accuracy: 0.6432 - loss: 0.9510
Epoch 2: val_accuracy improved from 0.70500 to 0.73250, saving model to resnet_model.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered
576/576 549s 953ms/step - accuracy: 0.6433 - loss: 0.9510 - val_accuracy: 0.7325 - val_loss: 0.6699
Epoch 3/15
576/576 0s 906ms/step - accuracy: 0.6830 - loss: 0.8185
Epoch 3: val_accuracy improved from 0.73250 to 0.75000, saving model to resnet_model.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered
576/576 558s 946ms/step - accuracy: 0.6830 - loss: 0.8185 - val_accuracy: 0.7500 - val_loss: 0.6317
Epoch 4/15
576/576 0s 905ms/step - accuracy: 0.7023 - loss: 0.7689
Epoch 4: val_accuracy improved from 0.75000 to 0.77750, saving model to resnet_model.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered
576/576 544s 944ms/step - accuracy: 0.7023 - loss: 0.7689 - val_accuracy: 0.7775 - val_loss: 0.5887
Epoch 5/15
576/576 0s 908ms/step - accuracy: 0.7159 - loss: 0.7173
Epoch 5: val_accuracy improved from 0.77750 to 0.78250, saving model to resnet_model.h5
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered
576/576 565s 949ms/step - accuracy: 0.7159 - loss: 0.7173 - val_accuracy: 0.7825 - val_loss: 0.5741
Epoch 6/15
576/576 0s 899ms/step - accuracy: 0.7275 - loss: 0.7106
Epoch 6: val_accuracy did not improve from 0.78250
576/576 556s 939ms/step - accuracy: 0.7275 - loss: 0.7105 - val_accuracy: 0.7350 - val_loss: 0.6148
Epoch 7/15
576/576 0s 901ms/step - accuracy: 0.7347 - loss: 0.6490
Epoch 7: val_accuracy did not improve from 0.78250
576/576 542s 941ms/step - accuracy: 0.7347 - loss: 0.6490 - val_accuracy: 0.7600 - val_loss: 0.5823
Epoch 8/15
576/576 0s 896ms/step - accuracy: 0.7406 - loss: 0.6502
Epoch 8: val_accuracy did not improve from 0.78250
576/576 557s 968ms/step - accuracy: 0.7406 - loss: 0.6502 - val_accuracy: 0.7725 - val_loss: 0.5727
Epoch 9/15
576/576 0s 899ms/step - accuracy: 0.7556 - loss: 0.6139
Epoch 9: val_accuracy did not improve from 0.78250
576/576 559s 970ms/step - accuracy: 0.7556 - loss: 0.6139 - val_accuracy: 0.7775 - val_loss: 0.5521
Epoch 10/15
576/576 0s 902ms/step - accuracy: 0.7629 - loss: 0.5986
Epoch 10: val_accuracy did not improve from 0.78250
576/576 542s 940ms/step - accuracy: 0.7629 - loss: 0.5986 - val_accuracy: 0.7750 - val_loss: 0.5827
Epoch 10: early stopping
```

```
import matplotlib.pyplot as plt

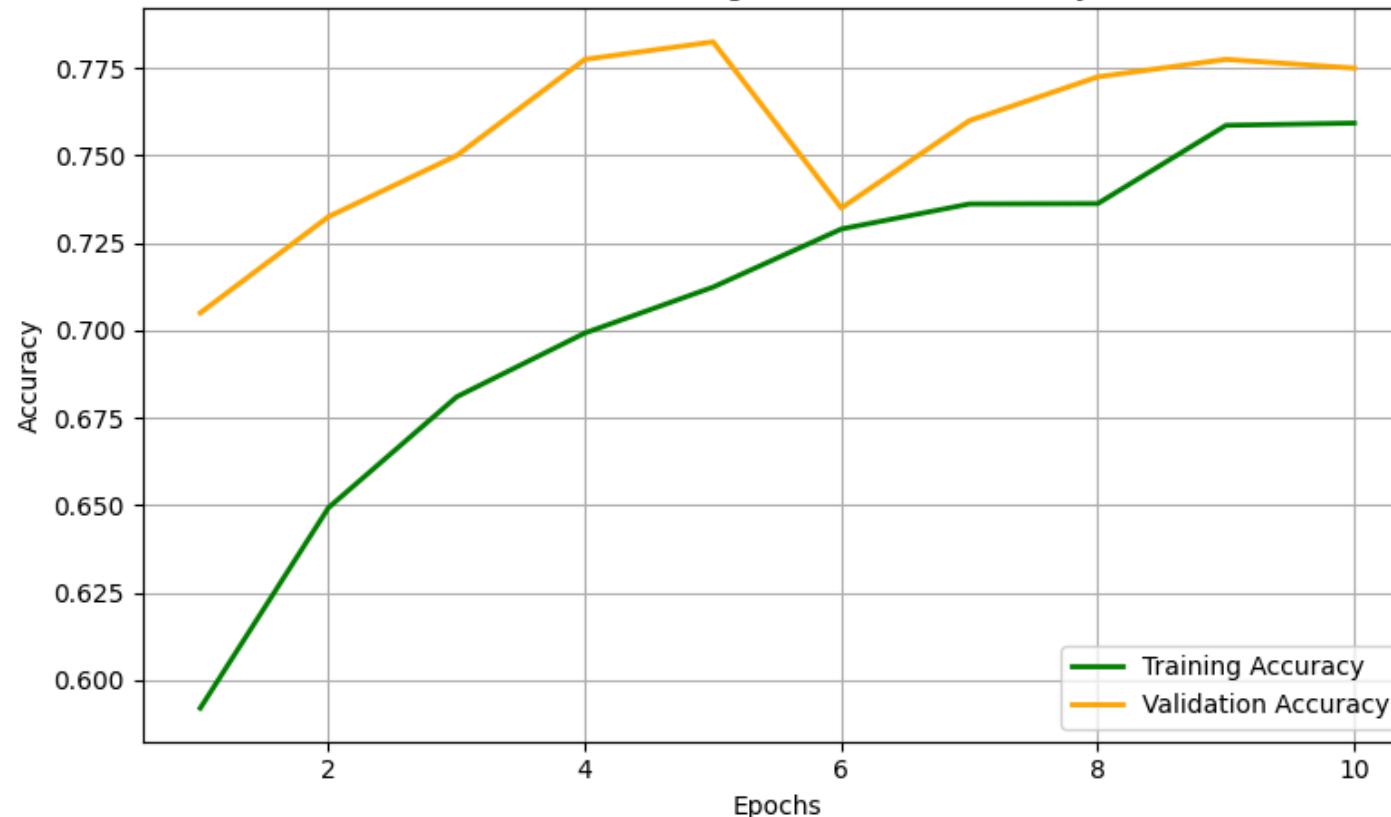
train_acc = resnet_history.history['accuracy']
val_acc = resnet_history.history['val_accuracy']
train_loss = resnet_history.history['loss']
val_loss = resnet_history.history['val_loss']
epochs = range(1, len(train_acc) + 1)

plt.figure(figsize=(8, 5))
plt.plot(epochs, train_acc, 'green', label='Training Accuracy', linewidth=2)
plt.plot(epochs, val_acc, 'orange', label='Validation Accuracy', linewidth=2)
plt.title('ResNet50: Training vs Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

plt.figure(figsize=(8, 5))
plt.plot(epochs, train_loss, 'red', label='Training Loss', linewidth=2)
plt.plot(epochs, val_loss, 'blue', label='Validation Loss', linewidth=2)
plt.title('ResNet50: Training vs Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right')
plt.grid(True)
plt.tight_layout()
plt.show()
```

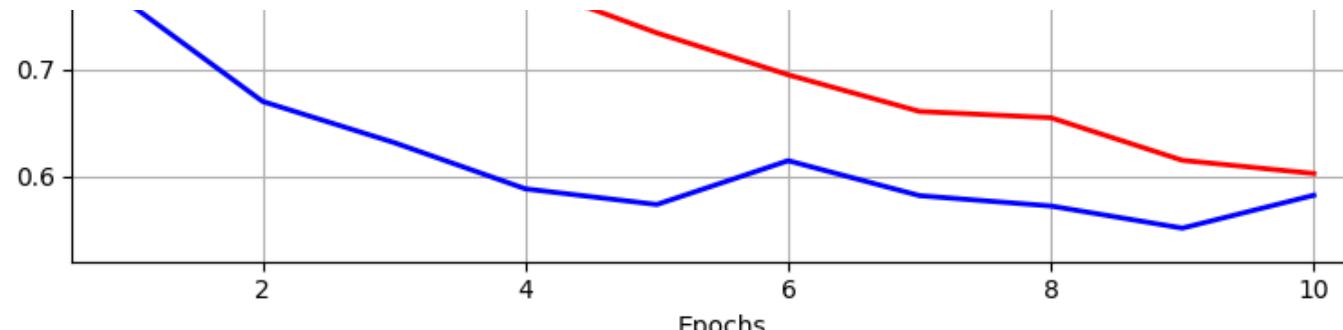


ResNet50: Training vs Validation Accuracy



ResNet50: Training vs Validation Loss





```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

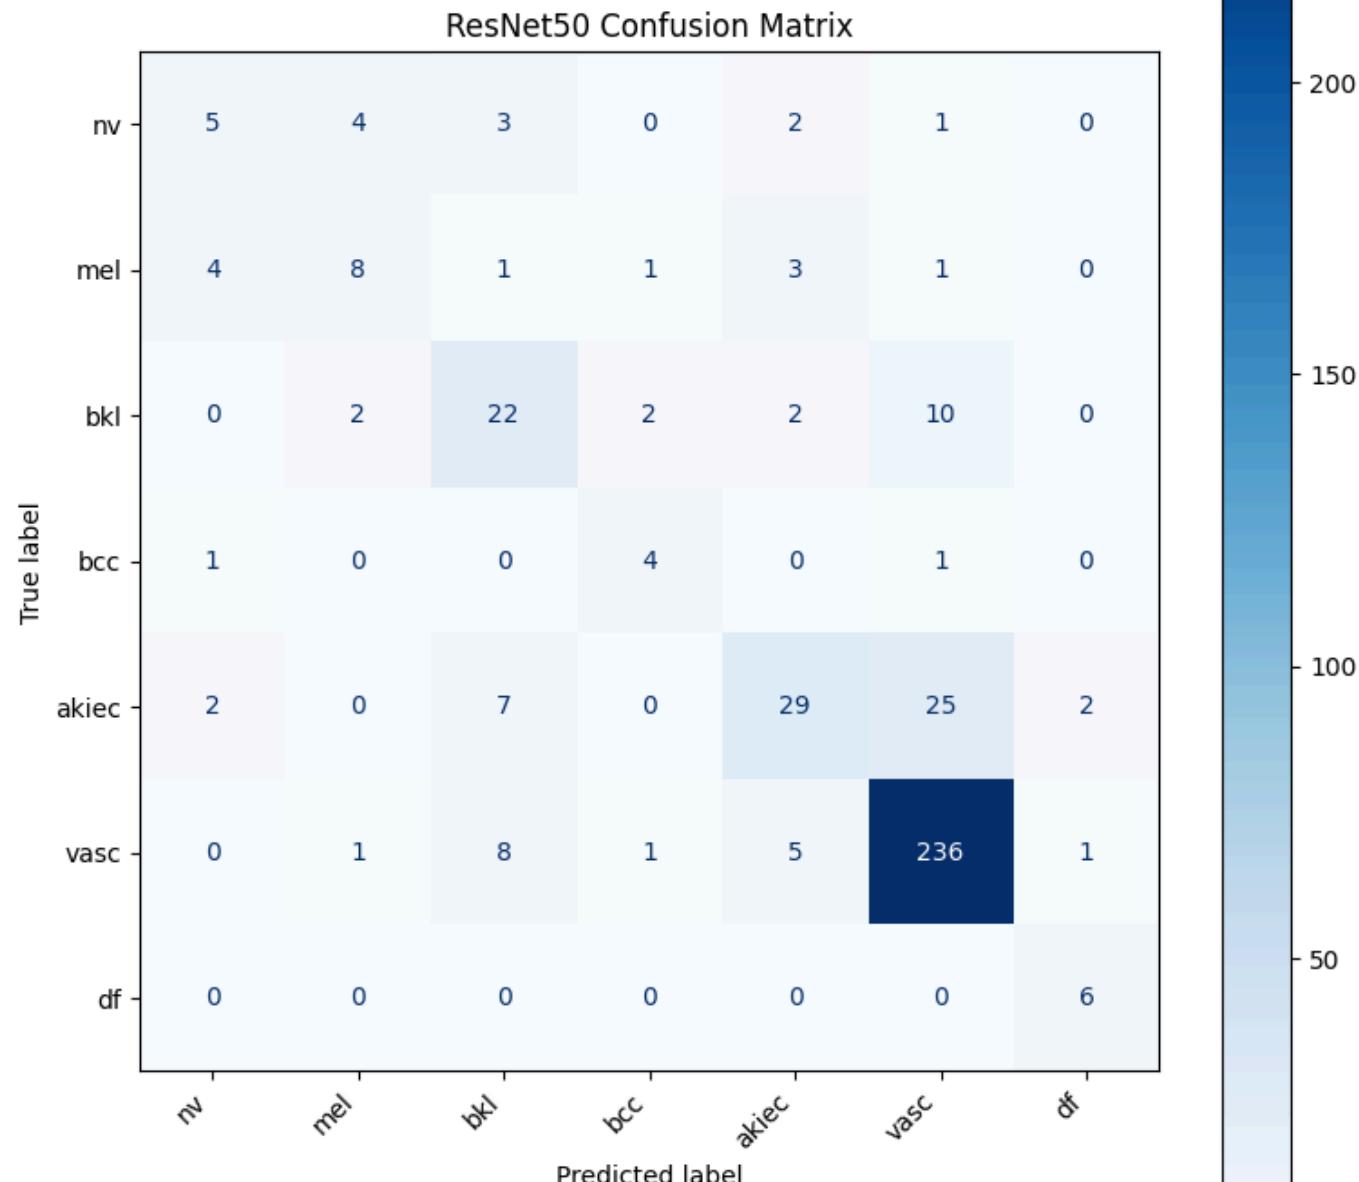
test_ds = val_ds
y_probs = resnet_model.predict(test_ds)
y_pred = np.argmax(y_probs, axis=1)

y_true_list = []
for X_batch, y_batch in test_ds:
    if y_batch.ndim > 1:
        y_true_list.append(np.argmax(y_batch, axis=1))
    else:
        y_true_list.append(y_batch)
y_true = np.concatenate(y_true_list, axis=0)

cm = confusion_matrix(y_true, y_pred)

disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                               display_labels=class_names)
fig, ax = plt.subplots(figsize=(8, 8))
disp.plot(ax=ax, cmap='Blues', colorbar=True)
plt.title('ResNet50 Confusion Matrix')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

25/25 25s 857ms/step



```
yolo_model = YOLO('/content/drive/MyDrive/yolov8m.pt')
data_yaml = f"""
train: {os.path.join(restructured_dir, 'train', 'images')}
val: {os.path.join(restructured_dir, 'val', 'images')}
nc: {len(class_names)}
names: {class_names}
"""

with open(os.path.join(restructured_dir, 'data.yaml'), 'w') as f:
    f.write(data_yaml)
yolo_model.train(data=os.path.join(restructured_dir, 'data.yaml'), epochs=3, imgsz=64, batch=32)
```

→ Ultralytics 8.3.114 🚀 Python-3.11.12 torch-2.6.0+cu124 CPU (Intel Xeon 2.20GHz)
engine/trainer: task=detect, mode=train, model=/content/drive/MyDrive/yolov8m.pt, data=/content/ham10000_restructured/data.yaml, epochs=3, Downloading <https://ultralytics.com/assets/Arial.ttf> to '/root/.config/Ultralytics/Arial.ttf'...
100% [██████████] 755k/755k [00:00<00:00, 10.6MB/s]Overriding model.yaml nc=80 with nc=7

	from	n	params	module	arguments
0		-1	1	1392 ultralytics.nn.modules.conv.Conv	[3, 48, 3, 2]
1		-1	1	41664 ultralytics.nn.modules.conv.Conv	[48, 96, 3, 2]
2		-1	2	111360 ultralytics.nn.modules.block.C2f	[96, 96, 2, True]
3		-1	1	166272 ultralytics.nn.modules.conv.Conv	[96, 192, 3, 2]
4		-1	4	813312 ultralytics.nn.modules.block.C2f	[192, 192, 4, True]
5		-1	1	664320 ultralytics.nn.modules.conv.Conv	[192, 384, 3, 2]
6		-1	4	3248640 ultralytics.nn.modules.block.C2f	[384, 384, 4, True]
7		-1	1	1991808 ultralytics.nn.modules.conv.Conv	[384, 576, 3, 2]
8		-1	2	3985920 ultralytics.nn.modules.block.C2f	[576, 576, 2, True]
9		-1	1	831168 ultralytics.nn.modules.block.SPPF	[576, 576, 5]
10		-1	1	0 torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
11		[-1, 6]	1	0 ultralytics.nn.modules.conv.Concat	[1]
12		-1	2	1993728 ultralytics.nn.modules.block.C2f	[960, 384, 2]
13		-1	1	0 torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
14		[-1, 4]	1	0 ultralytics.nn.modules.conv.Concat	[1]
15		-1	2	517632 ultralytics.nn.modules.block.C2f	[576, 192, 2]
16		-1	1	332160 ultralytics.nn.modules.conv.Conv	[192, 192, 3, 2]
17		[-1, 12]	1	0 ultralytics.nn.modules.conv.Concat	[1]
18		-1	2	1846272 ultralytics.nn.modules.block.C2f	[576, 384, 2]
19		-1	1	1327872 ultralytics.nn.modules.conv.Conv	[384, 384, 3, 2]
20		[-1, 9]	1	0 ultralytics.nn.modules.conv.Concat	[1]
21		-1	2	4207104 ultralytics.nn.modules.block.C2f	[960, 576, 2]
22		[15, 18, 21]	1	3779749 ultralytics.nn.modules.head.Detect	[7, [192, 384, 576]]

Model summary: 169 layers, 25,860,373 parameters, 25,860,357 gradients, 79.1 GFLOPs

Transferred 469/475 items from pretrained weights

Freezing layer 'model.22.dfl.conv.weight'

train: Fast image access (ping: 0.1±0.0 ms, read: 170.3±300.2 MB/s, size: 217.8 KB)

train: Scanning /content/ham10000_restructured/train/labels/akiec... 8012 images, 3500 backgrounds, 0 corrupt: 100% |██████████| 11512/11512

train: New cache created: /content/ham10000_restructured/train/labels/akiec.cache

albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, num_output_channels=3, method='weigh

val: Fast image access (ping: 0.0±0.0 ms, read: 73.1±7.9 MB/s, size: 288.9 KB)

val: Scanning /content/ham10000_restructured/val/labels/akiec... 2003 images, 0 backgrounds, 0 corrupt: 100% |██████████| 2003/2003 [00:05<0

val: New cache created: /content/ham10000_restructured/val/labels/akiec.cache

Plotting labels to runs/detect/train/labels.jpg...

optimizer: 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automati

optimizer: AdamW(lr=0.000909, momentum=0.9) with parameter groups 77 weight(decay=0.0), 84 weight(decay=0.0005), 83 bias(decay=0.0)

Image sizes 64 train, 64 val

Using 0 dataloader workers

Logging results to **runs/detect/train**

Starting training for 3 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/3	0G	0.4688	1.596	1.019	59	64: 100% ██████████ 360/360 [17:26<00:00, 2.91s/it]
	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 32/32 [01:21<00:00, 2.56s/it]

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
2/3	0G	0.2255	1.051	0.9189	60	64: 100% ██████████ 360/360 [16:54<00:00, 2.82s/it]

```
print(yolo_model)
```

```
YOLO(
  (model): DetectionModel(
    (model): Sequential(
      (0): Conv(
        (conv): Conv2d(3, 48, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
        (bn): BatchNorm2d(48, eps=0.001, momentum=0.03, affine=True, track_running_stats=True)
        (act): SiLU(inplace=True)
      )
      (1): Conv(
        (conv): Conv2d(48, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
        (bn): BatchNorm2d(96, eps=0.001, momentum=0.03, affine=True, track_running_stats=True)
        (act): SiLU(inplace=True)
      )
      (2): C2f(
        (cv1): Conv(
```

```
(conv): Conv2d(96, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn): BatchNorm2d(96, eps=0.001, momentum=0.03, affine=True, track_running_stats=True)
(act): SiLU(inplace=True)
)
(cv2): Conv(
    (conv): Conv2d(192, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn): BatchNorm2d(96, eps=0.001, momentum=0.03, affine=True, track_running_stats=True)
    (act): SiLU(inplace=True)
)
(m): ModuleList(
    (0-1): 2 x Bottleneck(
        (cv1): Conv(
            (conv): Conv2d(48, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn): BatchNorm2d(48, eps=0.001, momentum=0.03, affine=True, track_running_stats=True)
            (act): SiLU(inplace=True)
        )
        (cv2): Conv(
            (conv): Conv2d(48, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn): BatchNorm2d(48, eps=0.001, momentum=0.03, affine=True, track_running_stats=True)
            (act): SiLU(inplace=True)
        )
    )
)
)
(3): Conv(
    (conv): Conv2d(96, 192, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn): BatchNorm2d(192, eps=0.001, momentum=0.03, affine=True, track_running_stats=True)
    (act): SiLU(inplace=True)
)
(4): C2f(
    (cv1): Conv(
        (conv): Conv2d(192, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.03, affine=True, track_running_stats=True)
        (act): SiLU(inplace=True)
    )
    (cv2): Conv(
        (conv): Conv2d(576, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn): BatchNorm2d(192, eps=0.001, momentum=0.03, affine=True, track_running_stats=True)
        (act): SiLU(inplace=True)
    )
)
(m): ModuleList(
    (0-3): 4 x Bottleneck(
        (cv1): Conv(
```

```

def evaluate_classification(model, val_ds, model_name):
    val_labels = np.concatenate([y for _, y in val_ds], axis=0)
    predictions = model.predict(val_ds)
    y_true = np.argmax(val_labels, axis=1)
    y_pred = np.argmax(predictions, axis=1)
    print(f"\n{model_name} Classification Report:")
    print(classification_report(y_true, y_pred, target_names=[lesion_type_dict[c]['name'] for c in class_names]))
    severities = [lesion_type_dict[class_names[p]]['severity'] for p in y_pred]
    print(f"{model_name} Severity Distribution:")
    print(pd.Series(severities).value_counts())

def evaluate_yolo(model, val_dir, metadata):
    val_images_dir = os.path.join(val_dir, 'images')
    print(f"Looking for images in: {val_images_dir}")
    val_images = glob.glob(os.path.join(val_images_dir, '*', '*.jpg'))
    print(f"Found {len(val_images)} images in subdirectories")
    if not val_images:
        raise FileNotFoundError(f"No images found in {val_images_dir} or its subdirectories.")
    results = model.predict(source=val_images, save=False, imgs2=128)
    y_true, y_pred, sizes, severities = [], [], [], []
    pixel_to_mm = 0.1
    for result in results:
        image_id = os.path.basename(result.path).split('.')[0]
        true_label = metadata[metadata['image_id'] == image_id]['dx'].values[0]
        y_true.append(class_names.index(true_label))
        if len(result.bboxes) > 0:
            pred_class = int(result.bboxes.cls[0])
            confidence = float(result.bboxes.conf[0])
            y_pred.append(pred_class)
            box = result.bboxes.xywh[0]
            width, height = float(box[2]), float(box[3])
            size_mm2 = (width * pixel_to_mm) * (height * pixel_to_mm)
            sizes.append(size_mm2)
            severity = lesion_type_dict[class_names[pred_class]]['severity'] if confidence > 0.5 else 'Low'
            severities.append(severity)
        else:
            y_pred.append(-1)
            sizes.append(0)
            severities.append('Low')
    print("\nYOLOv8 Classification Report:")
    print(classification_report(y_true, [p if p != -1 else y_true[i] for i, p in enumerate(y_pred)], target_names=class_names))

```

```

print(f"Average Size of Infected Area: {np.mean(sizes):.2f} mm²")
print("YOLOv11 Severity Distribution:")
print(pd.Series(severities).value_counts())

val_images = glob.glob(os.path.join(restructured_dir, 'val', 'images', '*', '*.jpg'))
print(f"Number of validation images found: {len(val_images)}")
print(f"Validation metadata size: {len(val_meta)}")
print(val_meta.head())

```

→ Number of validation images found: 2003
 Validation metadata size: 2003

	lesion_id	image_id	dx	dx_type	age	sex	\
9725	HAM_0004376	ISIC_0024843	akiec	histo	70.0	female	
6059	HAM_0003024	ISIC_0024768	nv	follow_up	35.0	female	
4540	HAM_0001659	ISIC_0026564	nv	follow_up	35.0	male	
3817	HAM_0004625	ISIC_0029346	nv	follow_up	40.0	male	
7914	HAM_0000443	ISIC_0034271	nv	histo	35.0	female	

	localization	cell_type	severity	cell_type_idx
9725	face	Actinic keratoses	Moderate	0
6059	trunk	Melanocytic nevi	Low	4
4540	lower extremity	Melanocytic nevi	Low	4
3817	upper extremity	Melanocytic nevi	Low	4
7914	back	Melanocytic nevi	Low	4

```

evaluate_classification(cnn_model, val_ds, "CNN")
evaluate_classification(resnet_model, val_ds, "ResNet")

```

→ 25/25 ————— 5s 186ms/step

CNN Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Melanocytic nevi	0.46	0.40	0.43	15
Melanoma	0.33	0.44	0.38	18
Benign keratosis-like lesions	0.39	0.34	0.37	38
Basal cell carcinoma	0.30	0.50	0.38	6
Actinic keratoses	0.52	0.54	0.53	65
Vascular lesions	0.88	0.86	0.87	252
Dermatofibroma	0.67	0.67	0.67	6

accuracy	0.71	400
----------	------	-----

macro avg	0.51	0.54	0.52	400
weighted avg	0.72	0.71	0.72	400

CNN Severity Distribution:

Low 266

Moderate 100

High 34

Name: count, dtype: int64

25/25 ————— 26s 1s/step

ResNet Classification Report:

	precision	recall	f1-score	support
Melanocytic nevi	0.42	0.33	0.37	15
Melanoma	0.53	0.44	0.48	18
Benign keratosis-like lesions	0.54	0.58	0.56	38
Basal cell carcinoma	0.50	0.67	0.57	6
Actinic keratoses	0.71	0.45	0.55	65
Vascular lesions	0.86	0.94	0.90	252
Dermatofibroma	0.67	1.00	0.80	6
accuracy			0.78	400
macro avg	0.60	0.63	0.60	400
weighted avg	0.77	0.78	0.76	400

ResNet Severity Distribution:

Low 295

Moderate 82

High 23

Name: count, dtype: int64

```
def plot_history(history, title):
    plt.figure(figsize=(15, 5))
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Training Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title(f'{title} Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'], label='Training Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
```

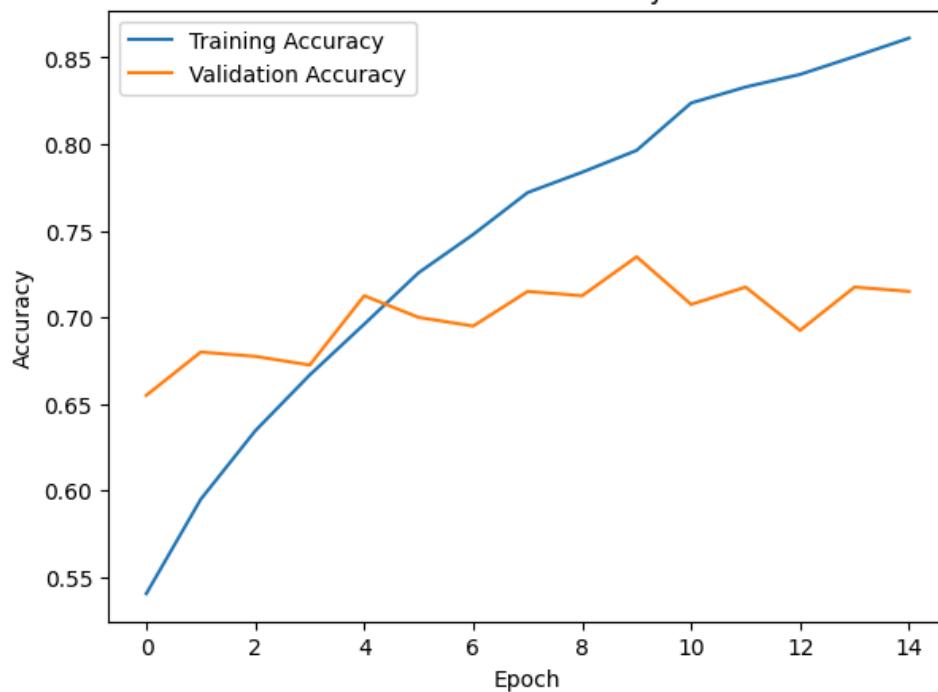
```
plt.title(f'{title} Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

def plot_yolo_predictions(model, val_dir, num_samples=5):
    val_images_dir = os.path.join(val_dir, 'images')
    val_images = glob.glob(os.path.join(val_images_dir, '*', '*.jpg'))
    print(f"Found {len(val_images)} images in subdirectories")
    if not val_images:
        raise FileNotFoundError(f"No images found in {val_images_dir} or its subdirectories.")
    results = model.predict(source=val_images, save=False, imgszt=128)
    plt.figure(figsize=(15, 10))
    for i, result in enumerate(results[:num_samples]):
        img = cv2.imread(result.path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        plt.subplot(1, num_samples, i+1)
        plt.imshow(img)
        if len(result.boxes) > 0:
            box = result.boxes.xyxy[0].numpy()
            cv2.rectangle(img, (int(box[0]), int(box[1])), (int(box[2]), int(box[3])), (255, 0, 0), 2)
            label = f"{class_names[int(result.boxes.cls[0])]} ({result.boxes.conf[0]:.2f})"
            plt.title(label)
        plt.axis('off')
    plt.show()

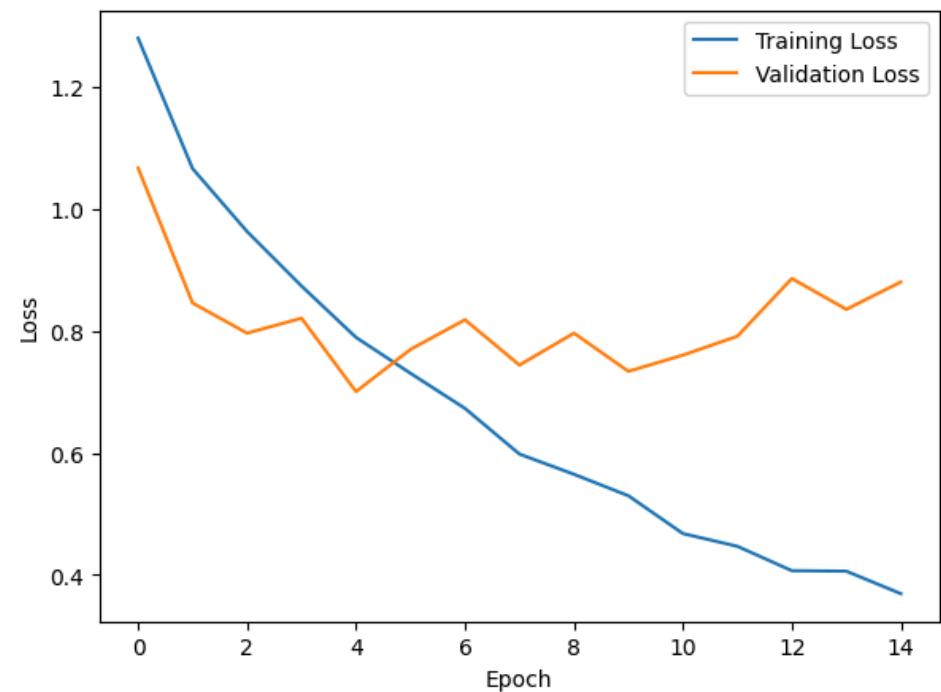
plot_history(cnn_history, "CNN Model")
plot_history(resnet_history, "ResNet Model")
```



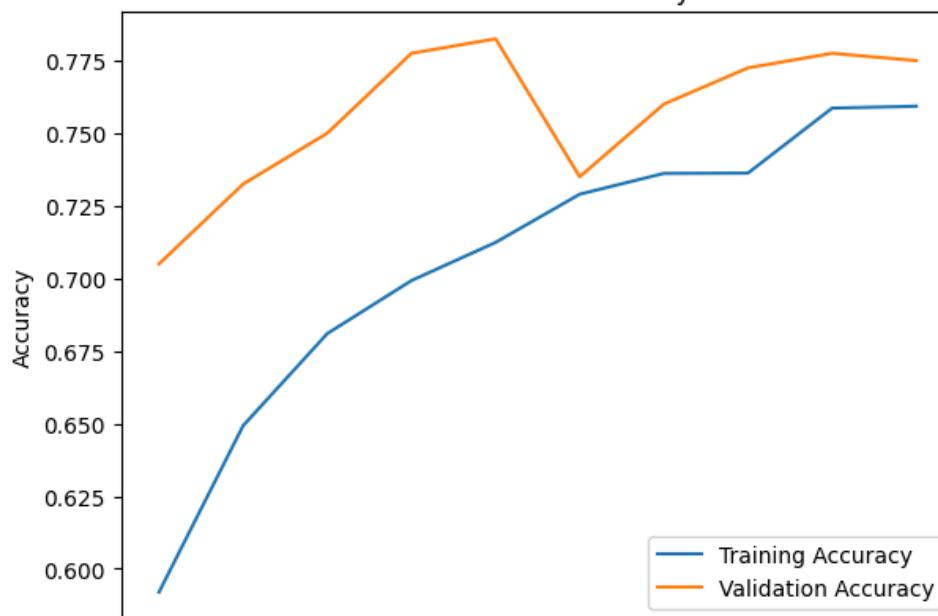
CNN Model Accuracy



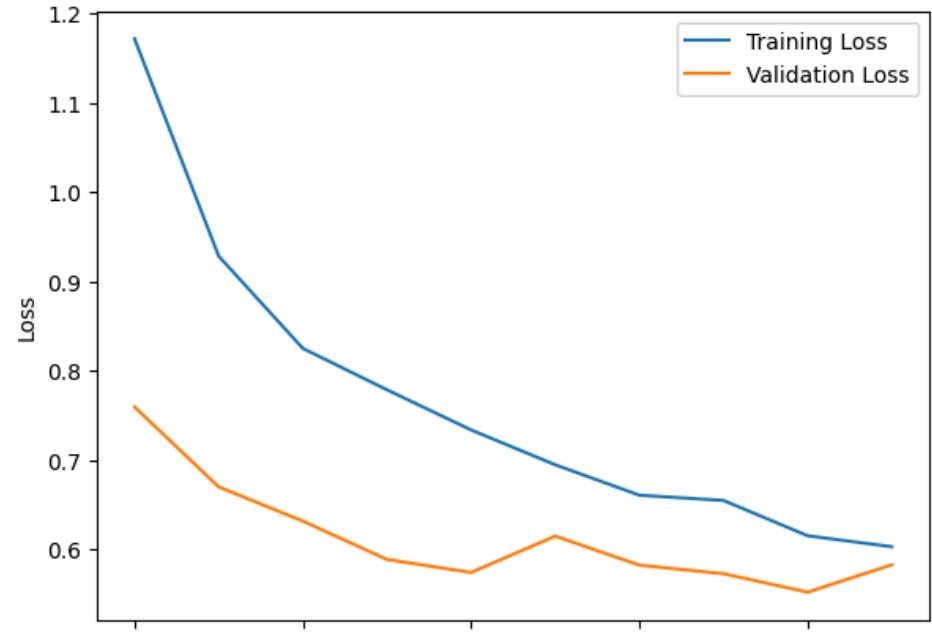
CNN Model Loss



ResNet Model Accuracy



ResNet Model Loss



```
from ultralytics import YOLO
from PIL import Image
import matplotlib.pyplot as plt
import os
import numpy as np

IMAGE_PATHS = [
    "/content/ham10000_restructured/val/images/nv/ISIC_0024321.jpg",
    "/content/ham10000_restructured/val/images/bkl/ISIC_0024381.jpg",
    "/content/ham10000_restructured/val/images/df/ISIC_0024330.jpg",
    "/content/ham10000_restructured/val/images/vasc/ISIC_0026456.jpg",
    "/content/ham10000_restructured/val/images/bkl/ISIC_0024358.jpg",
]
MODEL_PATH  = "/content/drive/MyDrive/best.pt"

custom_labels = {
    0: "Nevus",
    1: "Melanoma",
    2: "Benign Keratosis",
    3: "Basal Cell Carcinoma",
    4: "Actinic Keratoses",
    5: "Vascular Lesion",
    6: "Dermatofibroma"
}

folder_to_label = {
    "nv":    "Nevus",
    "mel":   "Melanoma",
    "bkl":   "Benign Keratosis",
    "bcc":   "Basal Cell Carcinoma",
    "akiec": "Actinic Keratoses",
    "vasc":  "Vascular Lesion",
```

```
"df": "Dermatofibroma"
}

model = YOLO(MODEL_PATH)

results = model(IMAGE_PATHS)

print("\n YOLO Skin Lesion Batch Results\n" + "-"*70)
for i, (path, res) in enumerate(zip(IMAGE_PATHS, results), start=1):
    actual_folder = os.path.basename(os.path.dirname(path))
    actual_label = folder_to_label.get(actual_folder, "Unknown")
    print(f"[Image {i}] {os.path.basename(path)}")
    print(f"  • Actual Class : {actual_label}")

    if res.boxes:
        for j, box in enumerate(res.boxes, start=1):
            cls_id = int(box.cls[0])
            conf = float(box.conf[0])
            pred = custom_labels.get(cls_id, f"Unknown({cls_id})")
            print(f"    Prediction {j}: {pred} (conf: {conf:.2f})")
    else:
        print("    No lesions detected.")
print("-"*70)

fig, axes = plt.subplots(1, 5, figsize=(20, 5))
for ax, res in zip(axes, results):
    annotated = res.plot()

    if annotated.shape[2] == 3 and annotated.dtype == np.uint8:
        annotated = annotated[..., ::-1]
    ax.imshow(annotated)
    ax.axis("off")
plt.suptitle("Annotated Skin Lesion Predictions", fontsize=16)
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```